

Solution of Robot Learning Written Part

ZHANG SHUTAO 1155238521

February 17, 2026

1 Problem 1: Inverted Pendulum Nonlinear Dynamics (Euler–Lagrange Derivation)

1.1 Model Setup and some assumptions

We consider a torque-actuated inverted pendulum with mass m , length l . And its point mass m at the end of a rigid, massless rod of length l . The inverted pendulum rotates around a fixed pivot. Let θ denote the generalized coordinate (pendulum angle), and let the control input be the applied torque u at the pivot. We also include viscous damping at the joint, modeled by a torque $-b\dot{\theta}$, where $b > 0$ is the damping coefficient.

1.2 Kinematics

Place the pivot at the origin and describe the position of the mass in Cartesian coordinates as:

$$x = l \sin \theta, \quad y = l \cos \theta \quad (1)$$

Taking time derivatives gives:

$$\dot{x} = l \cos \theta, \dot{\theta}, \quad \dot{y} = -l \sin \theta, \dot{\theta}. \quad (2)$$

Hence the squared speed is:

$$v^2 = \dot{x}^2 + \dot{y}^2 = l^2 \dot{\theta}^2 (\cos^2 \theta + \sin^2 \theta) = l^2 \dot{\theta}^2. \quad (3)$$

1.3 Kinetic Energy and Potential Energy

The kinetic energy can be written as:

$$T = \frac{1}{2}mv^2 = \frac{1}{2}ml^2\dot{\theta}^2 \quad (4)$$

Define the (effective) moment of inertia about the pivot:

$$I = ml^2, \quad (5)$$

so that

$$T = \frac{1}{2}I\dot{\theta}^2. \quad (6)$$

The gravitational potential energy can be written as:

$$V = mgy = mgl \cos \theta, \quad (7)$$

where g is the gravitational acceleration. (Any constant offset in V does not affect the equations of motion.)

1.4 Euler–Lagrange Equation with Generalized Force

Firstly, the Lagrangian is:

$$\mathcal{L}(\theta, \dot{\theta}) = T - V = \frac{1}{2}I\dot{\theta}^2 - mgl \cos \theta. \quad (8)$$

And for generalized coordinate θ , the Euler–Lagrange equation is

$$\frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{\theta}} \right) - \frac{\partial \mathcal{L}}{\partial \theta} = Q_\theta, \quad (9)$$

where Q_θ is the generalized torque applied along θ . Here we include the control torque u and viscous damping torque $-b\dot{\theta}$:

$$Q_\theta = u - b\dot{\theta}. \quad (10)$$

And expand the equation 9, compute each term of the equation 9:

$$\frac{\partial \mathcal{L}}{\partial \dot{\theta}} = \frac{\partial}{\partial \dot{\theta}} \left(\frac{1}{2}I\dot{\theta}^2 - mgl \cos \theta \right) = I\ddot{\theta}, \quad (11)$$

$$\frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{\theta}} \right) = \frac{d}{dt}(I\ddot{\theta}) = I\ddot{\theta}, \quad (12)$$

$$\frac{\partial \mathcal{L}}{\partial \theta} = \frac{\partial}{\partial \theta} \left(\frac{1}{2}I\dot{\theta}^2 - mgl \cos \theta \right) = mgl \sin \theta. \quad (13)$$

Substituting into the Euler–Lagrange equation we can yield:

$$I\ddot{\theta} - mgl \sin \theta = u - b\dot{\theta}. \quad (14)$$

Reordering terms gives the nonlinear equation of motion:

$$I\ddot{\theta} + b\dot{\theta} + mgl \sin \theta = u, \quad I = ml^2. \quad (15)$$

1.5 State-space Form

Define the state vector and the variables:

$$x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} \theta \\ \dot{\theta} \end{bmatrix}. \quad (16)$$

$$\dot{x}_1 = x_2, \quad (17)$$

And from the equation of motion:

$$\ddot{\theta} = \frac{1}{I} (u - b\dot{\theta} - mgl \sin \theta) = \frac{1}{I} (u - bx_2 - mgl \sin x_1). \quad (18)$$

Thus the nonlinear state-space dynamics are:

$$\dot{x} = \begin{bmatrix} \dot{x}_1 \\ \dot{x}_2 \end{bmatrix} = \begin{bmatrix} x_2 \\ \frac{1}{I} (u - bx_2 - mgl \sin x_1) \end{bmatrix}. \quad (19)$$

2 Problem 3: LQR (Optimal Control + Dynamic Programming)

2.1 Problem 3.1: Linearization and Discretization

2.1.1 Nonlinear Dynamics and State-space form

From problem 1, the inverted pendulum dynamics are given by:

$$I\ddot{\theta} + b\dot{\theta} + mgl \sin \theta = u, \quad I = ml^2, \quad (20)$$

where θ is the pendulum angle, u is the applied angle, b is viscous damping and mgl is the gravity term.

Define the state vector:

$$x \triangleq \begin{bmatrix} \theta \\ \dot{\theta} \end{bmatrix} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}. \quad (21)$$

Then the nonlinear dynamics system can be written as:

$$\dot{x} = f(x, u) = \begin{bmatrix} \dot{\theta} \\ \ddot{\theta} \end{bmatrix} = \begin{bmatrix} x_2 \\ \frac{1}{I} (u - bx_2 - mgl \sin x_1) \end{bmatrix}. \quad (22)$$

2.1.2 Equilibrium point for the upright position

We linearize around the upright equilibrium in a specified state:

$$\theta^* = \pi, \quad \dot{\theta}^* = 0, \quad u^* = 0. \quad (23)$$

This is an equilibrium because substituting $(\theta^*, \dot{\theta}^*, u^*)$ into (20) yields

$$I \cdot 0 + b \cdot 0 + mgl \sin(\pi) = 0 \Rightarrow 0 = 0. \quad (24)$$

2.1.3 Deviation (error) coordinates

For LQR, it is standard to work with deviation variables (small perturbations around the equilibrium):

$$\delta\theta \triangleq \theta - \pi, \quad \dot{\delta\theta} \triangleq \dot{\theta} - 0 = \dot{\theta}, \quad \delta u \triangleq u - 0 = u. \quad (25)$$

Collect them into:

$$\delta x \triangleq x - x^* = \begin{bmatrix} \delta\theta \\ \dot{\delta\theta} \end{bmatrix}. \quad (26)$$

In the remainder, we abuse notation and simply write x to mean the deviation state δx , since the linearized dynamics are expressed in deviation coordinates.

2.1.4 Linearization method A (using the $\sin(\pi + \delta\theta) \approx -\delta\theta$)

The hint suggests using:

$$\sin(\pi + \delta\theta) \approx -\delta\theta \quad \text{for small } \delta\theta. \quad (27)$$

Substitute $\theta = \pi + \delta\theta$ into (20):

$$I\ddot{\theta} + b\dot{\theta} + mgl \sin(\pi + \delta\theta) = u. \quad (28)$$

Using (27), we get

$$I\ddot{\theta} + b\dot{\theta} - mgl \delta\theta = u. \quad (29)$$

Because $\theta = \pi + \delta\theta$, we have $\dot{\theta} = \delta\dot{\theta}$ and $\ddot{\theta} = \delta\ddot{\theta}$, hence

$$I\delta\ddot{\theta} + b\delta\dot{\theta} - mgl \delta\theta = \delta u. \quad (30)$$

Solve for $\delta\ddot{\theta}$:

$$\delta\ddot{\theta} = \frac{mgl}{I}\delta\theta - \frac{b}{I}\delta\dot{\theta} + \frac{1}{I}\delta u. \quad (31)$$

Now write the first-order state-space form with $x = [\delta\theta, \dot{\delta\theta}]^\top$:

$$\dot{x} = \begin{bmatrix} \dot{\delta\theta} \\ \ddot{\delta\theta} \end{bmatrix} = \begin{bmatrix} x_2 \\ \frac{mgl}{I}x_1 - \frac{b}{I}x_2 + \frac{1}{I}u \end{bmatrix}. \quad (32)$$

Therefore the continuous-time linearized dynamics are:

$$\dot{x} = Ax + Bu, \quad (33)$$

with

$$A = \begin{bmatrix} 0 & 1 \\ \frac{mgl}{I} & -\frac{b}{I} \end{bmatrix}, \quad B = \begin{bmatrix} 0 \\ \frac{1}{I} \end{bmatrix}. \quad (34)$$

2.1.5 Linearization method B (Jacobian Method)

As a consistency check, linearization can also be obtained via Jacobian:

$$\delta \dot{x} \approx \frac{\partial f}{\partial x} \Big|_{(x^*, u^*)} \delta x + \frac{\partial f}{\partial u} \Big|_{(x^*, u^*)} \delta u. \quad (35)$$

From (22), we have:

$$f(x, u) = \begin{bmatrix} x_2 \\ \frac{1}{I}(u - bx_2 - mgl \sin x_1) \end{bmatrix}. \quad (36)$$

After which, we can compute the partial derivatives:

$$\frac{\partial f}{\partial x} = \begin{bmatrix} \frac{\partial x_2}{\partial x_1} & \frac{\partial x_2}{\partial x_2} \\ \frac{\partial}{\partial x_1} \left(\frac{u - bx_2 - mgl \sin x_1}{I} \right) & \frac{\partial}{\partial x_2} \left(\frac{u - bx_2 - mgl \sin x_1}{I} \right) \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ -\frac{mgl}{I} \cos x_1 & -\frac{b}{I} \end{bmatrix}. \quad (37)$$

At $x_1 = \theta^* = \pi$, $\cos \pi = -1$, so

$$\frac{\partial f}{\partial x} \Big|_{(x^*, u^*)} = \begin{bmatrix} 0 & 1 \\ \frac{mgl}{I} & -\frac{b}{I} \end{bmatrix} = A,$$

and

$$\frac{\partial f}{\partial u} \Big|_{(x^*, u^*)} = \begin{bmatrix} 0 \\ \frac{1}{I} \end{bmatrix} = B.$$

This matches (34).

2.1.6 Euler discretization

Euler discretization is:

$$x_{k+1} = x_k + \Delta t (Ax_k + Bu_k). \quad (38)$$

Rearranging:

$$x_{k+1} = (I + \Delta t A)x_k + (\Delta t B)u_k. \quad (39)$$

Thus the discrete-time matrices are

$$A_d = I + \Delta t A, \quad B_d = \Delta t B. \quad (40)$$

Substituting (34), we obtain explicit forms:

$$A_d = \begin{bmatrix} 1 & \Delta t \\ \Delta t \frac{mgl}{I} & 1 - \Delta t \frac{b}{I} \end{bmatrix}, \quad B_d = \begin{bmatrix} 0 \\ \frac{\Delta t}{I} \end{bmatrix}. \quad (41)$$

2.2 Problem 3.2: Finite-horizon Discrete-time LQR via Dynamic Programming

2.2.1 Problem statement

Consider the discrete-time linear system

$$x_{k+1} = A_d x_k + B_d u_k, \quad k = 0, 1, \dots, N-1, \quad (42)$$

and the finite-horizon quadratic cost

$$J = \sum_{k=0}^{N-1} (x_k^\top Q x_k + u_k^\top R u_k) + x_N^\top Q_f x_N, \quad (43)$$

where $Q \succeq 0$, $R \succ 0$, and $Q_f \succeq 0$. The goal is to find an optimal sequence $\{u_k^*\}_{k=0}^{N-1}$ minimizing J .

2.2.2 Cost-to-go and Bellman optimality

Define the optimal cost-to-go at time k :

$$J_k^*(x) \triangleq \min_{\{u_i\}_{i=k}^{N-1}} \left[\sum_{i=k}^{N-1} (x_i^\top Q x_i + u_i^\top R u_i) + x_N^\top Q_f x_N \right], \quad (44)$$

with terminal condition

$$J_N^*(x) = x^\top Q_f x. \quad (45)$$

Bellman's optimality principle yields the recursion

$$J_k^*(x) = \min_u (x^\top Q x + u^\top R u + J_{k+1}^*(x^+)), \quad x^+ = A_d x + B_d u, \quad (46)$$

where x denotes the current state x_k and x^+ denotes the next state x_{k+1} .

2.2.3 Induction hypothesis: quadratic form of the value function

We claim $J_k^*(x)$ is a quadratic form:

$$J_k^*(x) = x^\top P_k x, \quad (47)$$

for some symmetric matrix $P_k \succeq 0$.

Base case. At $k = N$, (45) implies $P_N = Q_f$.

Inductive step. Assume (47) holds for $k+1$, i.e.,

$$J_{k+1}^*(x) = x^\top P_{k+1} x. \quad (48)$$

We show it then holds for k and derive the recursion for P_k .

2.2.4 Substitute the quadratic hypothesis into Bellman

Substitute (48) into (46):

$$J_k^*(x) = \min_u (x^\top Qx + u^\top Rx + (A_dx + B_du)^\top P_{k+1}(A_dx + B_du)). \quad (49)$$

2.2.5 Expand the quadratic expression

Expand the last term:

$$\begin{aligned} (A_dx + B_du)^\top P_{k+1}(A_dx + B_du) &= x^\top A_d^\top P_{k+1} A_dx + x^\top A_d^\top P_{k+1} B_d u + u^\top B_d^\top P_{k+1} A_dx + u^\top B_d^\top P_{k+1} B_d u \\ &= x^\top A_d^\top P_{k+1} A_dx + 2u^\top B_d^\top P_{k+1} A_dx + u^\top B_d^\top P_{k+1} B_d u, \end{aligned} \quad (50)$$

where the middle two terms are transposes of each other (scalars) and thus sum to twice one of them.

Plug (50) into (49):

$$J_k^*(x) = \min_u \left[x^\top (Q + A_d^\top P_{k+1} A_d)x + 2u^\top (B_d^\top P_{k+1} A_d)x + u^\top (R + B_d^\top P_{k+1} B_d)u \right]. \quad (51)$$

For convenience, define

$$S_k \triangleq R + B_d^\top P_{k+1} B_d, \quad G_k \triangleq B_d^\top P_{k+1} A_d, \quad H_k \triangleq Q + A_d^\top P_{k+1} A_d. \quad (52)$$

Then the minimized expression becomes

$$\Phi(u; x) = x^\top H_k x + 2u^\top G_k x + u^\top S_k u, \quad J_k^*(x) = \min_u \Phi(u; x). \quad (53)$$

2.2.6 Minimize with respect to u (first-order optimality condition)

Because $R \succ 0$ and $B_d^\top P_{k+1} B_d \succeq 0$, we have

$$S_k = R + B_d^\top P_{k+1} B_d \succ 0, \quad (54)$$

so $\Phi(u; x)$ is a strictly convex quadratic function of u , hence the minimizer is unique and can be found by setting the gradient to zero.

Compute the gradient w.r.t. u :

$$\nabla_u(u^\top S_k u) = 2S_k u, \quad \nabla_u(2u^\top G_k x) = 2G_k x, \quad \nabla_u(x^\top H_k x) = 0.$$

Thus:

$$\nabla_u \Phi(u; x) = 2S_k u + 2G_k x = 0 \implies S_k u^* = -G_k x. \quad (55)$$

Therefore,

$$u^* = -S_k^{-1} G_k x. \quad (56)$$

Define the time-varying feedback gain

$$K_k \triangleq S_k^{-1} G_k = (R + B_d^\top P_{k+1} B_d)^{-1} B_d^\top P_{k+1} A_d, \quad (57)$$

so that the optimal control law is

$$u_k^* = -K_k x_k. \quad (58)$$

2.2.7 Substitute the optimal u^* back to obtain the Riccati recursion

To obtain P_k , we compute the minimized value $\min_u \Phi(u; x)$. A clean way is to “complete the square”:

$$\begin{aligned}\Phi(u; x) &= x^\top H_k x + 2u^\top G_k x + u^\top S_k u \\ &= x^\top H_k x + (u + S_k^{-1} G_k x)^\top S_k (u + S_k^{-1} G_k x) - x^\top G_k^\top S_k^{-1} G_k x.\end{aligned}\quad (59)$$

The middle term is always ≥ 0 and is minimized to 0 at $u^* = -S_k^{-1} G_k x$. Hence,

$$J_k^*(x) = \min_u \Phi(u; x) = x^\top (H_k - G_k^\top S_k^{-1} G_k) x. \quad (60)$$

Therefore $J_k^*(x)$ is quadratic with

$$P_k = H_k - G_k^\top S_k^{-1} G_k. \quad (61)$$

Substitute definitions (52) into (61):

$$P_k = Q + A_d^\top P_{k+1} A_d - A_d^\top P_{k+1} B_d (R + B_d^\top P_{k+1} B_d)^{-1} B_d^\top P_{k+1} A_d. \quad (62)$$

Together with the terminal condition

$$P_N = Q_f, \quad (63)$$

we obtain the finite-horizon backward Riccati recursion.

2.2.8 Final boxed result (what we implement)

The finite-horizon discrete-time LQR solution is:

$$\begin{aligned}P_N &= Q_f, \\ K_k &= (R + B_d^\top P_{k+1} B_d)^{-1} B_d^\top P_{k+1} A_d, \\ u_k^* &= -K_k x_k, \\ P_k &= Q + A_d^\top P_{k+1} A_d - A_d^\top P_{k+1} B_d (R + B_d^\top P_{k+1} B_d)^{-1} B_d^\top P_{k+1} A_d, \quad k = N-1, \dots, 0.\end{aligned}\quad (64)$$

3 Problem 3.4 Conceptual Questions

3.1 Increasing Q typically:

Answer: (a)

Explanation: Increasing Q penalizes state error more, leading to stronger control actions and faster convergence.

3.2 Increasing R typically:

Answer: (b)

Explanation: Increasing R penalizes control effort more, resulting in smaller torques and slower system response.

3.3 Increasing Q_f typically:

Answer: (a)

Explanation: Increasing Q_f places more weight on the final state, emphasizing terminal accuracy.

ROSE 5760 Robot Learning Assignment 1


Coding Results Part

ZHANG SHUTAO 1155238521

Problem 1.3: Free Dynamics Tests

1. Default arguments


Open-Loop Inverted Pendulum Model Test ($u = 0$), $x_0 = [3.54, 0.00]$




2. Initial condition: $[\theta_0, \dot{\theta}_0] = [0, 0]$ and $[0 + 0.2, 0.3]$

Open-Loop Inverted Pendulum Model Test ($u = 0$), $x_0 = [0.00, 0.00]$

Open-Loop Inverted Pendulum Model Test ($u = 0$), $x_0 = [0.20, 0.30]$




3. Initial condition: $[\theta_0, \dot{\theta}_0] = [\pi, 0]$ and $[\pi, +0.001, 0]$




Problem 2: PID Control

1. PID Evaluation



Problem 3: Linear Quadratic Regulator (LQR)


1. LQR Implementation




Problem 4: Supervised Learning from LQR


1. MLP Policy

```
(robole) zhang@ZHANG:~/zst/robot_learning$ python train_mlp.py
Epoch 0001 | train 0.307837 | val 0.303708
Epoch 0005 | train 0.260436 | val 0.276873
Epoch 0010 | train 0.234905 | val 0.252749
Epoch 0015 | train 0.180592 | val 0.194910
Epoch 0020 | train 0.165138 | val 0.140903
Epoch 0025 | train 0.056354 | val 0.063179
Epoch 0030 | train 0.024322 | val 0.024593
Epoch 0035 | train 0.010686 | val 0.014245
Epoch 0040 | train 0.004305 | val 0.002952
Epoch 0045 | train 0.001779 | val 0.002353
Epoch 0050 | train 0.005509 | val 0.001207
Saved: models/mlp.pt
```





2. CNN Policy



3. LSTM Policy

